Arpit Parihar

Github | LinkedIn | arpitparihar07@gmail.com | San Jose, CA | +1.773.837.8826

Work Experience

Staff AI Engineer, Snowflake Inc. Feb 2022 - Apr 2025, Nov 2025 - Present

- Extending Snowflake's AI usage agent solo from usage reporting into effectiveness measurement, piloting with ∼130 engineers: scores suggestion acceptance and PR throughput, mines telemetry for high-leverage usage patterns, and serves them as inline recommendations via a Cortex Code plugin
- Built and lead the evaluation harness for internal and Cortex AI agents, piloted on the Snowflake sales agent, the highest-usage agent at the company, and currently onboarding the remaining GTM agents:
  - Scored agent trajectories, tool calls, correctness of generated SQL, and response quality via golden datasets, LLM-as-judge, and human labels
  - Added pulse-check evals against third-party MCP servers to catch upstream breakage before users did
  - Gated releases with regression runs in CI, catching 18 tool outages and 53 stale eval cases in the first week
  - Built the pilot solo, now leading a team of 2 extending coverage across agents
- Developed an advanced AI-based employee assistant chatbot for internal Snowflake documentation and knowledge base articles, currently used by 9000 employees for day-to-day inquiries:
  - Implemented a ColBERT v2 late-interaction retrieval system with contextualized embeddings and token-level MaxSim scoring, improving retrieval accuracy 40% over traditional approaches at sub-second query latency
  - Generated responses with Claude 3.5 through Snowflake Cortex, and built conversational agents that book meeting rooms and file IT tickets
  - Led development of the frontend, with citations, chat history, feedback, and custom bot creation

Applied Scientist, Delphina Inc. Apr 2025 - Nov 2025

- Owned modeling capabilities for Delphina's machine learning agent (team of 8), driving toward a "vibe ML" workflow where the user states an objective in plain text and the agent runs the modeling loop:
  - Broadened regression coverage with Poisson and negative binomial link functions for count data, and extended the agent past classification and regression by adapting its tree-based ensembles to time-series forecasting via column offsetting, later adding FFT features
  - Inserted an interpreter agent into the modeling loop to read each run's results and set the next run's hyperparameter space and tuning budget, alongside the existing learning-to-rank feature selection, reporting its reasoning per run
- After the company pivoted to an analytics agent, designed its fan-out subagent architecture for long-form recurring reporting (WBRs, QBRs) spanning many metrics:
  - Split the work across planning, metric deep-dive, synthesizer, and review subagents with strict contracts between them to limit information loss in LLM-to-LLM handoff
  - Built correctness and style evals in Braintrust, reaching a 85% pass rate across 5 design partners, one of which moved its WBR process onto Delphina entirely

Teaching Assistant, University of Chicago Mar 2021 - Oct 2024

- Instructed classes of up to 60 students in Machine Learning, Python for ML, Reinforcement Learning, Linear & Non-Linear Models, and MLOps, plus one-on-one sessions.

Data Science Intern, McKinsey & Company Jun 2021 - Sep 2021

- Built ensemble models predicting churn among farmers and distributors at 74% accuracy, run as experiments in MLFlow and Azure Databricks.

Senior Data Scientist, Boston Consulting Group Apr 2018 - May 2019

- Led the development of a deep learning model to detect breast cancer from mammograms, achieving an accuracy of 96% using Keras and convolutional neural networks.
- Built a Gradient Boosting model for a leading car seat manufacturer, deployed as an enterprise-scale real-time predictive maintenance framework in Python and Azure, scoring an F1 of 76% and saving $450k per plant yearly.

Data Science Manager, Flipkart Sep 2017 - Mar 2018

- Led a team siting optimal delivery hubs for a logistics company in Karnataka using geocoding and hierarchical clustering, cutting average delivery times by 50 minutes.

Data Scientist, Mu Sigma Mar 2015 - Aug 2017

- Flagged suspicious health insurance claims with a rule-based fraud engine, then trained a random forest classifier to prioritize them, saving $11M.

Education

2020 - 2021	Master's in Machine Learning at **University of Chicago**	(GPA: 4.0/4.0)
2010 - 2014	Bachelor's in Civil Engineering at **College of Engineering, Pune**	(GPA: 7.5/10.0)